Deep learning for query performance prediction reading list

- Sriraman, A., & Dhanotia, A. (2020, March). Accelerometer: Understanding Acceleration Opportunities for Data Center Overheads at Hyperscale. In *Proceedings of the Twenty-Fifth International Conference on Architectural Support for Programming Languages and Operating Systems* (pp. 733-750).

- Sun, J., & Li, G. (2019). An end-to-end learning-based cost estimator. *Proceedings of the VLDB Endowment*, *13*(3), 307-319.

- Holze, M., Haschimi, A., & Ritter, N. (2010, March). Towards workload-aware self-management: Predicting significant workload shifts. In *2010 IEEE 26th International Conference on Data Engineering Workshops (ICDEW 2010)* (pp. 111-116). IEEE.

- Holze, M., Gaidies, C., & Ritter, N. (2009, November). Consistent on-line classification of DBS workload events. In *Proceedings of the 18th ACM conference on Information and knowledge management* (pp. 1641-1644).

- Kraska, T., Alizadeh, M., Beutel, A., Chi, E. H., Ding, J., Kristo, A., ... & Nathan, V. (2019). Sagedb: A learned database system.

- Pavlo, A., Angulo, G., Arulraj, J., Lin, H., Lin, J., Ma, L., ... & Santurkar, S. (2017, January). Self-Driving Database Management Systems. In *CIDR* (Vol. 4, p. 1).

- Dutt, A., Wang, C., Nazi, A., Kandula, S., Narasayya, V., & Chaudhuri, S. (2019). Selectivity estimation for range predicates using lightweight models. *Proceedings of the VLDB Endowment*, *12*(9), 1044-1057.

- Mou, L., Li, G., Zhang, L., Wang, T., & Jin, Z. (2016, February). Convolutional neural networks over tree structures for programming language processing. In *Thirtieth AAAI Conference on Artificial Intelligence*.

- Mansouri, B., Rohatgi, S., Oard, D. W., Wu, J., Giles, C. L., & Zanibbi, R. (2019, September). Tangent-cft: An embedding model for mathematical formulas. In *Proceedings of the 2019 ACM SIGIR International Conference on Theory of Information Retrieval* (pp. 11-18).

- Mao, H., Schwarzkopf, M., Venkatakrishnan, S. B., Meng, Z., & Alizadeh, M. (2019). Learning scheduling algorithms for data processing clusters. In *Proceedings of the ACM Special Interest Group on Data Communication* (pp. 270-288).

- Ma, L., Van Aken, D., Hefny, A., Mezerhane, G., Pavlo, A., & Gordon, G. J. (2018, May). Query-based workload forecasting for self-driving database management systems. In *Proceedings of the 2018 International Conference on Management of Data* (pp. 631-645).

- Lample, G., & Charton, F. (2019). Deep learning for symbolic mathematics. *arXiv preprint arXiv:1912.01412*.

- Marcus, R., & Papaemmanouil, O. (2019). Plan-structured deep neural network models for query performance prediction. *Proceedings of the VLDB Endowment*, *12*(11), 1733-1746.

- Stillger, M., Lohman, G. M., Markl, V., & Kandil, M. (2001, September). LEO-DB2's learning optimizer. In *VLDB* (Vol. 1, pp. 19-28).

- Marcus, R., Negi, P., Mao, H., Zhang, C., Alizadeh, M., Kraska, T., ... & Tatbul, N. (2019). Neo: A learned query optimizer. *Proceedings of the VLDB Endowment*, *12*(11), 1705-1718.

- Ortiz, J., Balazinska, M., Gehrke, J., & Keerthi, S. S. (2018, June). Learning state representations for query optimization with deep reinforcement learning. In *Proceedings of the Second Workshop on Data Management for End-To-End Machine Learning* (pp. 1-4).

- Jindal, A., Patel, H., Roy, A., Qiao, S., Yin, Z., Sen, R., & Krishnan, S. (2019, November). Peregrine: Workload Optimization for Cloud Query Engines. In *Proceedings of the ACM Symposium on Cloud Computing* (pp. 416-427).

- Sen, R., & Ramachandra, K. (2018, February). Characterizing resource sensitivity of database workloads. In *2018 IEEE International Symposium on High Performance Computer Architecture (HPCA)* (pp. 657-669). IEEE.

- Kul, G., Luong, D. T. A., Xie, T., Chandola, V., Kennedy, O., & Upadhyaya, S. (2018). Similarity metrics for sql query clustering. *IEEE Transactions on Knowledge and Data Engineering*, *30*(12), 2408-2420.

- Rajan, K., Kakadia, D., Curino, C., & Krishnan, S. (2016, October). PerfOrator: eloquent performance models for resource optimization. In *Proceedings of the Seventh ACM Symposium on Cloud Computing* (pp. 415-427)